

TAMANNA

[Github](#) | [LinkedIn](#) | [Leetcode](#)

tamannas10110@gmail.com | +91-9877553894

AI Engineer with 5+ years designing and delivering production-grade Gen AI applications across cloud and enterprise environments — from architecture and prototyping through deployment and production operations. Specialized in building agentic AI systems and retrieval-augmented generation (RAG) workflows for enterprise document management, including document ingestion, processing, classification, and retrieval, while maintaining responsible AI practices and data security/compliance standards. Proficient in Python and React, with strong grounding in Gen AI frameworks (LangChain, LangGraph), prompt and context engineering, and software engineering best practices across the development lifecycle. Experienced integrating AI platforms with enterprise SaaS ecosystems via secure API-driven architectures, and working independently with general guidance to deliver scalable, reliable solutions adopted across large enterprise user bases.

WORK EXPERIENCE

Genpact (Senior Full Stack Engineer)	Bangalore, Karnataka July 2024 - Present
• Owned end-to-end client/business onboarding for an agentic AI capability (Genpact AP Trace) — from requirement gathering and solution design through configuration, validation, and production go-live — embedding with cross-functional teams (finance, compliance, operations) to conduct technical discovery and map workflow requirements; shipped a multi-agent system (Python, LangChain/LangGraph) flagging duplicates across 10,000+ invoices/month, reducing financial leakage risk by 15–20%. • Designed custom workflow and integration logic connecting enterprise platforms (ServiceNow, Salesforce) with internal systems via authenticated REST APIs, building technical solutions that bridged gaps between platform capabilities and client/business system requirements. • Built and debugged production LLM-based applications — multi-agent orchestration workflows (LangGraph) with state management and tool-calling — for vendor onboarding and compliance validation (Azure OpenAI), troubleshooting and resolving issues across the stack to cut manual processing effort by 40%. • Developed POCs and prototypes for GenAI document-processing pipelines (Azure OpenAI, Azure Document Intelligence), iterating from prototype to production and achieving 90%+ extraction accuracy across 500+ samples; monitored output quality and triaged production failures post-launch. • Built and shipped a conversational, LLM-based Text-to-SQL application (Python, FastAPI, PostgreSQL) — a natural-language interface for non-technical users — adopted by 200+ users across 10+ regions, reducing manual SQL dependency by 60%. • Contributed to onboarding frameworks and reusable tooling — designing policy-driven orchestration patterns and configurable templates across regions — accelerating onboarding by ~35% and cutting rollout time from 6 weeks to under 4. • Communicated technical solutions across audiences — explaining production issues and configuration changes to both technical engineering teams and non-technical business stakeholders — using a structured, discovery-first problem-solving approach. • Stayed current with emerging Gen AI and LLM-based application techniques (Azure cloud, Python), evaluating opportunities to apply them to evolving client and business needs.	
CGI (Software Engineer)	Bangalore, Karnataka June 2021 – June 2024
• Owned backend services end-to-end — design, implementation, deployment, and operational health — for cloud-native enterprise applications (React.js, Python — Flask/FastAPI), supporting 200+ enterprise users. • Designed REST API architectures (FastAPI, PostgreSQL, Redis) with strong backend development fundamentals, focusing on reliability and performance — reducing average API response time by 35%. • Built Python-based automations and third-party integrations, converting recurring operational issues into engineered, maintainable workflow solutions that replaced manual, legacy enterprise processes — reducing processing effort by 40–50% and eliminating 15–20 hrs/week of manual work. • Developed reusable React component libraries and coding standards, raising code quality and consistency across the team's solutions — cutting feature delivery time by 30%. • Built Selenium-based automation scripts for critical enterprise workflows, reducing failures by 30%; maintained technical documentation, used source control, and participated in code reviews for collaborative development and debugging. • Operated as a fully independent practitioner across the full application lifecycle with minimal supervision — requirements, architecture, build, and deployment — applying strong debugging instincts across legacy systems and ambiguous environments. • Integrated with third-party APIs as part of internal automation tooling, using PostgreSQL and Redis for backend data storage and caching needs.	

EDUCATION

Qualification	Institute	CGPA / Percentage	Time of Completion
Bachelor of Technology	NIT Jalandhar	CGPA – 7.92/10	August 2017 – May 2021
Class XII	GGN Public School, Ludhiana	91.8 %	April 2016 – March 2017
Class X	GGN Public School, Ludhiana	CGPA – 10/10	April 2014 – March 2015

TECHNICAL SKILLS

- **AI / GenAI / Agentic AI:** LLMs, Generative AI, Agentic AI, Multi-Agent Orchestration, LangChain, LangGraph, RAG, Prompt Engineering, Azure OpenAI, Azure Document Intelligence, Vector Databases (FAISS/Pinecone/ChromaDB), Function/Tool Calling, Scikit-Learn, TensorFlow, LLM finetuning using PEFT, RLHF
- **Languages:** Python, JavaScript, SQL, C++
- **Frontend:** React.js, Redux Toolkit, MUI (Material-UI), Tailwind CSS, HTML5/CSS3
- **Backend & APIs:** FastAPI, Flask, REST APIs, WebSockets, Microservices Architecture
- **Databases & Caching:** PostgreSQL, Redis, MongoDB
- **Cloud & DevOps:** Azure (App Service, Functions, OpenAI Service), Docker, Git, CI/CD
- **Data & Visualization:** NumPy, Pandas, Matplotlib, Seaborn
- **Automation & Testing:** Selenium, PyTest
- **CS Fundamentals:** Data Structures & Algorithms, Object-Oriented Programming, System Design

PERSONAL PROJECT

Agentic Medical Appointments Processing – [DEMO](#)

The system simulates how medical appointments move through an AI-orchestrated pipeline:

- **Orchestrator** — A master agent reads appointments from a database and ranks them by dynamic priority (client, specialty, notes).
- **Stage agents** — Each appointment passes through 6 configurable stages (e.g. Intake, Insurance, Clinical Review). Each stage is an LLM-backed agent that returns a standardized outcome.
- **Exception queue** — If a stage returns Escalate, the case goes to a human concierge.
- **Human resolution** — A human resolves the issue; the appointment becomes Cleared and the workflow resumes automatically.

Tech Stack Used – React, Vite, FastAPI, Python, SQLite, LangChain, LangGraph, Hugging Face LLM

CERTIFICATIONS

- [Document AI: From OCR to Agentic Doc Extraction](#) – May 2026
- [Databricks Certified Generative AI Engineer Associate](#)– July 2025