

TAMANNA

[Github](#) | [LinkedIn](#) | [Leetcode](#)

tamannas10110@gmail.com | +91-9877553894

AI Engineer with 5+ years designing and delivering production-grade, cloud-native AI-native applications end-to-end - from architecture and prototyping through deployment and production operations. Specialized in building agentic AI systems and retrieval-augmented generation (RAG) workflows that integrate large language models into real enterprise processes, scaling across large user bases and multiple regions. Experienced connecting AI platforms with enterprise SaaS ecosystems through API-driven integrations, with strong grounding in prompt and context engineering, orchestration frameworks, and scalable, reliable application architecture. Combines this with full-stack engineering capability (React, Python) to translate business and product requirements into working, adopted solutions - embedding closely with cross-functional and business stakeholders to balance speed, quality, and long-term maintainability.

WORK EXPERIENCE

Genpact (Forward Deployed Application Engineer)	Bangalore, Karnataka July 2024 - Present
• Designed, developed, and operated production-grade, cloud-native enterprise applications integrating AI capabilities — large language models, retrieval-augmented generation (RAG), and agent-based workflows — for Genpact AP Trace, an agentic AI invoice-duplicate-detection platform; architected multi-agent logic flagging duplicates across 10,000+ invoices/month, reducing financial leakage risk by 15–20%. • Implemented RAG architectures and agentic systems to ground LLM outputs in enterprise document data, applying architectural guardrails and evaluation frameworks to ensure compliant, trustworthy, and high-quality AI outputs while reducing hallucination in production workflows. • Architected multi-agent orchestration frameworks (LangGraph, Azure OpenAI) to automate vendor onboarding and compliance validation - applying prompt engineering, context engineering, and responsible, effective use of AI agents within enterprise operations - cutting manual processing effort by 40%. • Owned application delivery end-to-end — from problem framing and prototyping through deployment and production operations - for a Text-to-SQL AI application; embedded with cross-functional and business stakeholders to rapidly understand operational workflows, translating requirements into a scalable, secure, and reliable solution adopted by 200+ users across 10+ regions, reducing manual SQL dependency by 60%. • Built GenAI-powered document intelligence pipelines (Azure OpenAI, Azure Document Intelligence) on cloud-native, SaaS infrastructure, achieving 90%+ extraction accuracy across 500+ samples while balancing speed, quality, and long-term maintainability. • Rapidly prototyped, iterated, and productionized reusable reference architectures and assets for region-specific business logic, adapting them to real-world enterprise constraints - accelerated onboarding by ~35%, cutting rollout time from 6 weeks to under 4. • Integrated enterprise SaaS platforms (ServiceNow, Salesforce) via REST API-driven architectures, strengthening scalability, security, and reliability; contributed reusable components and engineering best practices to improve efficiency and consistency across solutions. • Acted as a key technical interface to internal stakeholders and cross-functional teams, clearly communicating solution approaches and trade-offs; proactively managed delivery risks and technical escalations, ensuring milestones and quality standards were met within scope. • Demonstrated a proactive, learning-mindset approach to leveraging AI in everyday engineering workflows, exercising technical leadership and accountability in ambiguous, high-impact environments. • Operated across Linux and Azure cloud environments with hands-on observability and MLOps - adjacent practices - monitoring pipeline reliability, extraction accuracy, and AI output quality in production.	
CGI (Software Engineer)	Bangalore, Karnataka June 2021 – June 2024
• Owned the full application lifecycle - from design and implementation through deployment and production operations — for cloud-native enterprise applications (React.js, Python-Flask/FastAPI), supporting 200+ enterprise users; led solution architecture translating business process challenges into scalable, secure, and reliable software. • Designed REST API architectures (FastAPI, PostgreSQL, Redis) grounded in core architectural principles - scalability, security, and reliability - reducing average API response time by 35%. • Built Python-based automation pipelines, replacing manual, legacy enterprise processes with workflow automation comparable to iPaaS-style integration, reducing processing effort by 40-50% and eliminating 15-20 hrs/week of manual operational work.	

- **Developed reusable React component libraries, coding standards, and reusable assets/templates** to improve engineering efficiency and consistency across solutions, cutting feature delivery time by 30%.
- **Built Selenium-based process automation** for critical enterprise workflows, reducing failures by 30%; acted as a **technical interface to business stakeholders**, leading requirements gathering, demos, and technical documentation, clearly communicating solution approaches and trade-offs to technical and non-technical audiences.
- **Operated as a fully independent practitioner across the full application lifecycle** - requirements, architecture, build, and deployment - demonstrating **adaptability and a learning mindset in ambiguous, high-impact environments**, frequently navigating legacy/antiquated systems.
- Worked across **Linux and Windows environments**, using command-line tooling for builds, deployments, and debugging; contributed to **engineering standards and best practices**, maintaining strong accountability to internal teammates and stakeholders relying on delivered systems.
- **Collaborated cross-functionally with business stakeholders** to gather requirements and translate evolving operational needs into functional features - proactively managing delivery risks and technical trade-offs across multiple tools and technical domains.

EDUCATION

Qualification	Institute	CGPA / Percentage	Time of Completion
Bachelor of Technology	NIT Jalandhar	CGPA – 7.92/10	August 2017 – May 2021
Class XII	GGN Public School, Ludhiana	91.8 %	April 2016 – March 2017
Class X	GGN Public School, Ludhiana	CGPA – 10/10	April 2014 – March 2015

TECHNICAL SKILLS

- **AI / GenAI / Agentic AI:** LLMs, Generative AI, Agentic AI, Multi-Agent Orchestration, LangChain, LangGraph, RAG, Prompt Engineering, Azure OpenAI, Azure Document Intelligence, Vector Databases (FAISS/Pinecone/ChromaDB), Function/Tool Calling, Scikit-Learn, TensorFlow
- **Languages:** Python, JavaScript, SQL, C++
- **Frontend:** React.js, Redux Toolkit, MUI (Material-UI), Tailwind CSS, HTML5/CSS3
- **Backend & APIs:** FastAPI, Flask, REST APIs, WebSockets, Microservices Architecture
- **Databases & Caching:** PostgreSQL, Redis, MongoDB
- **Cloud & DevOps:** Azure (App Service, Functions, OpenAI Service), Docker, Git, CI/CD
- **Data & Visualization:** NumPy, Pandas, Matplotlib, Seaborn
- **Automation & Testing:** Selenium, PyTest
- **CS Fundamentals:** Data Structures & Algorithms, Object-Oriented Programming, System Design

PERSONAL PROJECT

Agentic Medical Appointments Processing – [DEMO](#)

The system simulates how medical appointments move through an AI-orchestrated pipeline:

- **Orchestrator** — A master agent reads appointments from a database and ranks them by dynamic priority (client, specialty, notes).
- **Stage agents** — Each appointment passes through 6 configurable stages (e.g. Intake, Insurance, Clinical Review). Each stage is an LLM-backed agent that returns a standardized outcome.
- **Exception queue** — If a stage returns Escalate, the case goes to a human concierge.
- **Human resolution** — A human resolves the issue; the appointment becomes Cleared and the workflow resumes automatically.

Tech Stack Used – React, Vite, FastAPI, Python, SQLite, LangChain, LangGraph, Hugging Face LLM

CERTIFICATIONS

- [Document AI: From OCR to Agentic Doc Extraction](#) – May 2026
- [Databricks Certified Generative AI Engineer Associate](#)– July 2025